

The goal of this document is to show how to develop JAVA microservices that can interact with Kafka, and, with a database available in a cloud environment. The project should be built from these indications.

The usage of a cloud database requires an Amazon AWS account.

The previous Tutorial “Deploying Kafka service in Amazon AWS EC2” is used as a service used by Kafka Producer and Kafka Consumer.

The following picture overviews the technological and application server dependencies that are introduced by this tutorial.

Description:

Any application service that requires the production or consumption of Kafka messages, in a predefined topic, invokes JVM and refers to the JAVA class and the required libraries, 1 and 2 are exemplified, more application services could be used. *Producer depends on org.apache.kafka libs and Consumer depends on org.apache.kafka libs and mysql.jdbc libs.*

In this example, both Producer and Consumer are configured to use the same Kafka server installed on EC2 Amazon AWS. After that, the Consumer propagates the received messages to an RDS Amazon database configured as a MySQL service. This architecture allows a light production and consumption of messages, and then storing the received messages for future use in a different repository.

Then, a Dreamfactory server is configured to deploy a REST interface that exposes the datamodel of the RDS Amazon database. All the following REST operations are allowed: GET, POST, PUT, PATCH, DELETE. These REST operations are consumed by the BPEL and BPMN models that are deployed in ORACLE SOA platform.

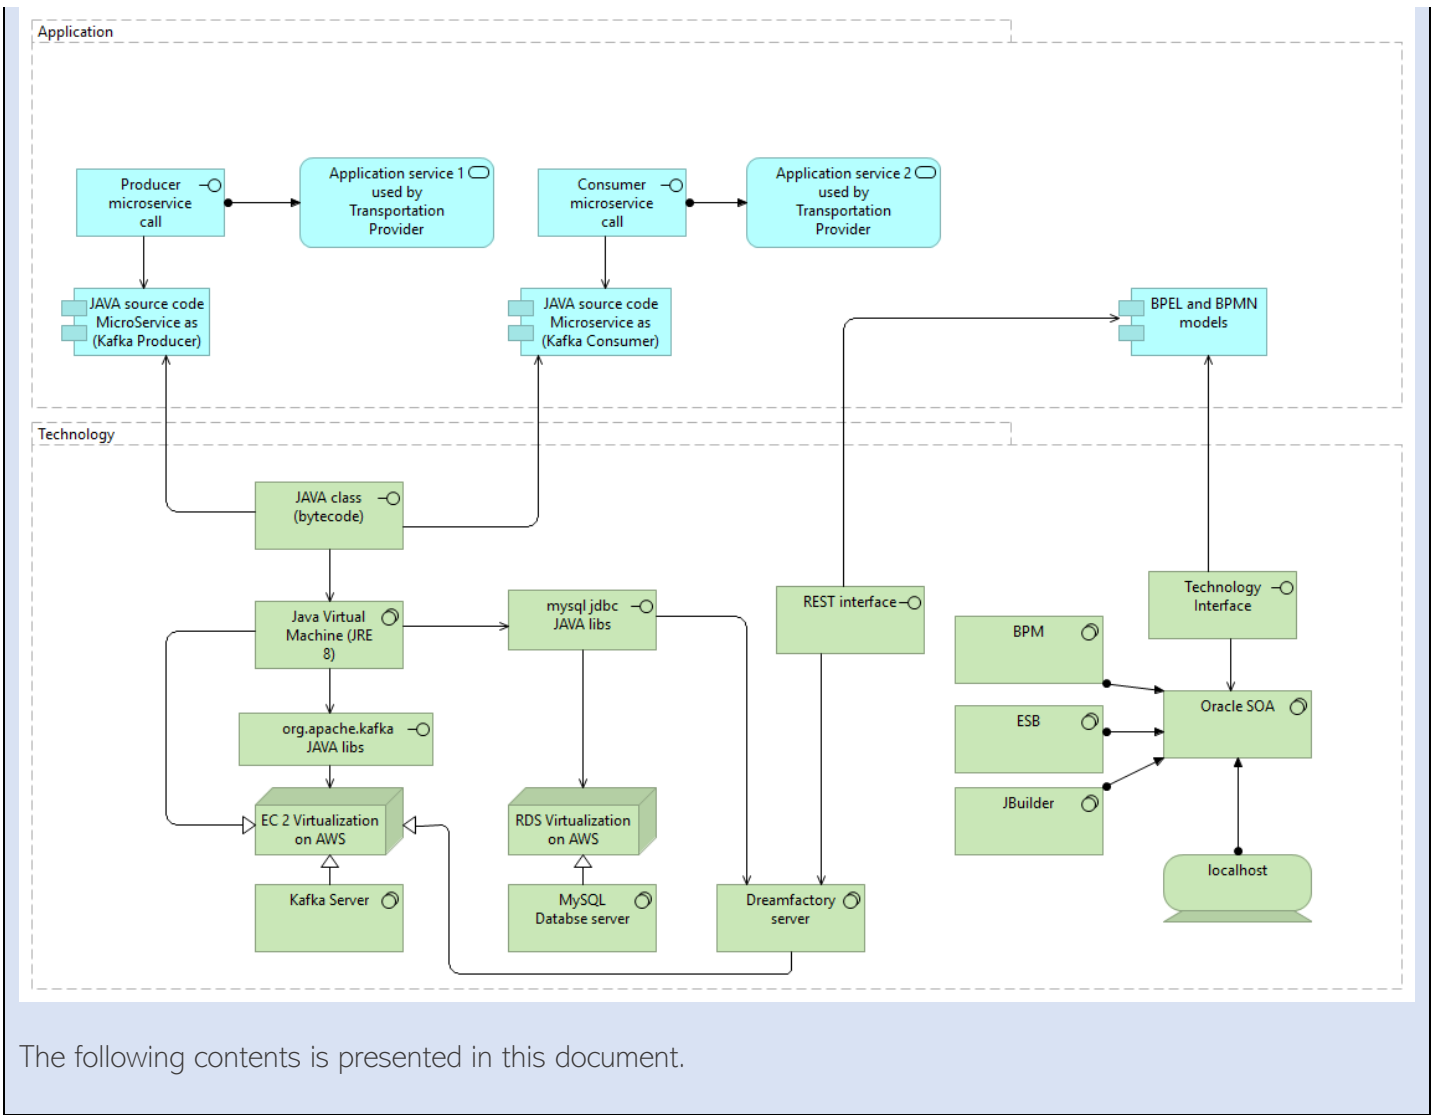
Scalability:

JAVA microservices could be developed and deployed on-the-fly accordingly with business needs.

EC2 virtualization could be partitioned and clustered if more performance is demanded.

RDS virtualization could be escalated, creating different SCHEMAS, or DATABASES instances, for different Application services.

A single DreamFactory server could be configured for multiple SCHEMAS, or DATABASES instances. Exposing each configuration as a different service with its own WEBAPP_id in the URL.



The following contents is presented in this document.

Contents

| | |
|---|----|
| A. Prerequisites for using Maven in Eclipse: | 3 |
| B. Create a JAVA Kafka Producer (for windows): example of | 4 |
| C. Create a JAVA Kafka Consumer (for windows): example of | 7 |
| D. Producing and Consuming an XML message (for windows): example of | 8 |
| E. Create an AWS RDS Database, accessing it, and inserting data: example of | 9 |
| F. Deploying JAVA services in EC2: example of | 12 |
| G. Dreamfactory installation and operation in AWS EC2 | 15 |
| H. Creating a topic in Kafka from JAVA | 19 |
| I. Appendix: URLs with other resources | 20 |

A. Prerequisites for using Maven in Eclipse:

- A.0. On the top menu bar, open Window -> Show View -> Other
- A.1. In the Show View window, open Maven -> Maven Repositories
- A.2. In the window that appears, right-click on Global Repositories and select Go Into
- A.3. Right-click on "central (http://repo.maven.apache.org/maven2)" and select "Rebuild Index"
- A.4. Note that it will take a while to complete the download
- A.5. Once indexing is complete, Right-click on the project -> Maven -> Add Dependency and start typing the name of the project you want to import (such as "hibernate").
- A.6. The search results will auto-fill in the "Search Results" box below.

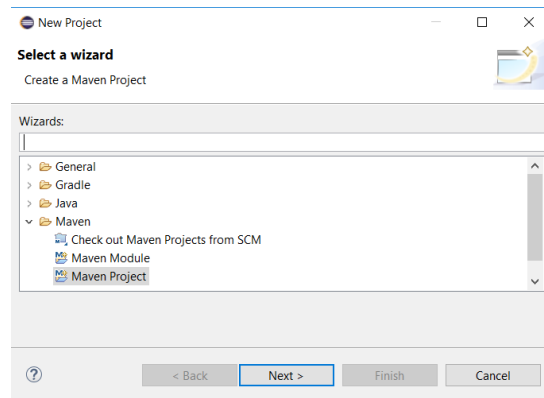
Then to install Maven plugin in Eclipse:

Help -> install new software -> <http://download.eclipse.org/technology/m2e/releases> Add

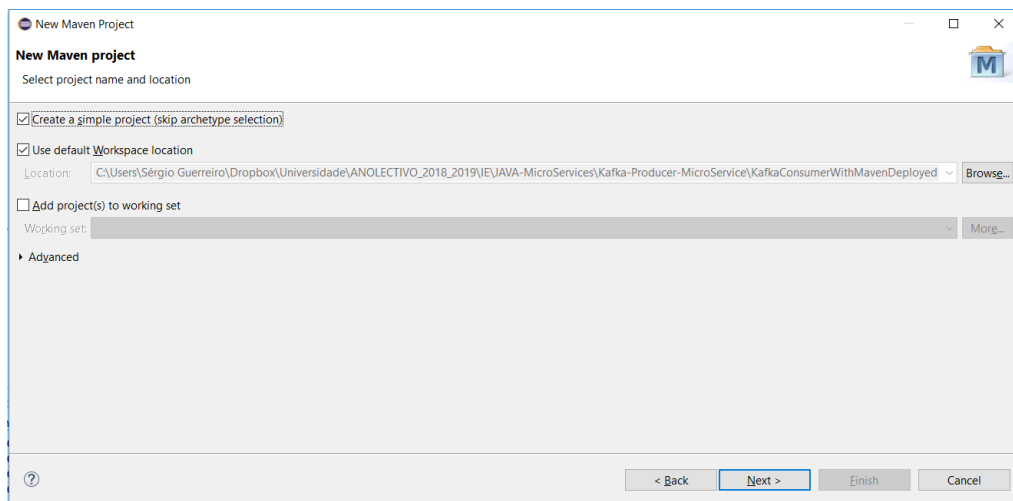
B. Create a JAVA Kafka Producer (for windows): example of

The goal of this section is to create a micro service in JAVA to produce messages for a Kafka service that has already been provisioned and a topic has been created.

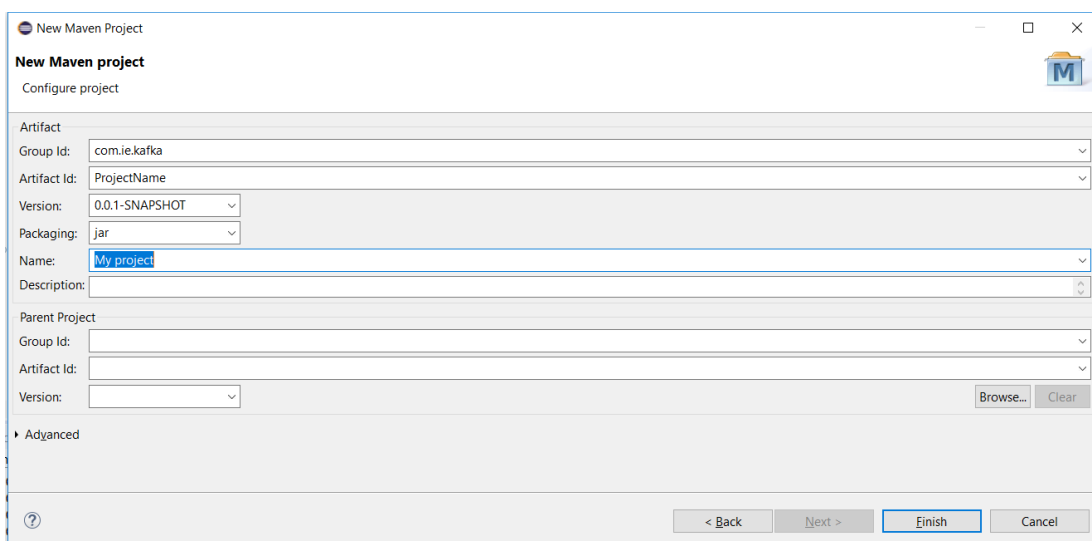
1. Create a Maven project. File > New > Project...



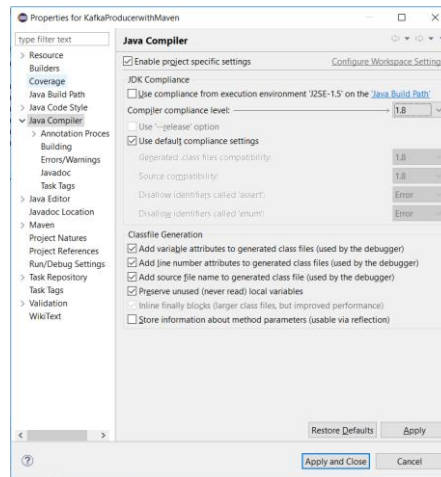
2. Check the option "Create a simple project (skip archetype selection)" and then press Next



Use the following configuration (consult Appendix: URLs with other resources further details). And then press Finish.



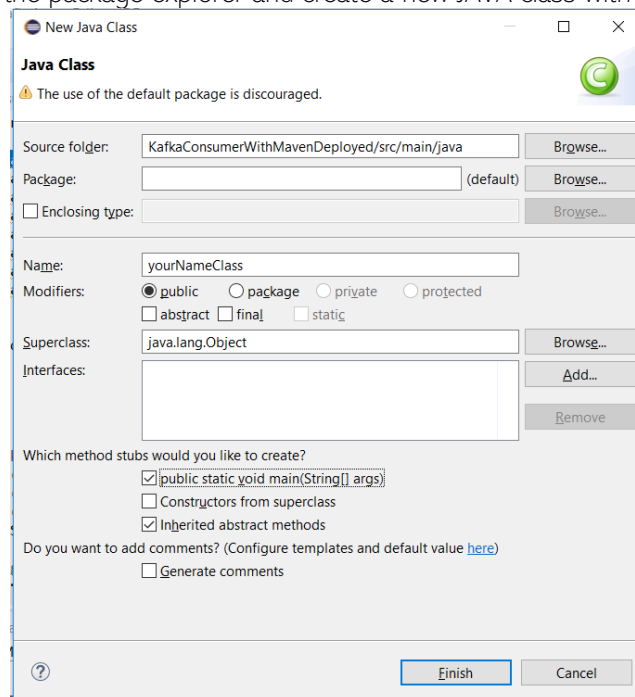
3. In Project > Properties > Java compiler check that compilation is done using java 8



4. Define the following Maven dependency in pom.xml of your project:

```
<dependencies>
<dependency>
  <groupId>org.apache.kafka</groupId>
  <artifactId>kafka-clients</artifactId>
  <version>2.1.0</version>
</dependency>
</dependencies>
```

5. Right click in your project in the package explorer and create a new JAVA class with the following configuration:



6. Study the following source code and adapt it to your new class:

```
import java.util.Properties;
import org.apache.kafka.clients.producer.Callback;
import org.apache.kafka.clients.producer.KafkaProducer;
import org.apache.kafka.clients.producer.Producer;
import org.apache.kafka.clients.producer.ProducerConfig;
import org.apache.kafka.clients.producer.ProducerRecord;
import org.apache.kafka.common.serialization.LongSerializer;
import org.apache.kafka.common.serialization.StringSerializer;

public class ClientProducerKafka {
```

```

    public static void main(String[] args) {
        Properties kafkaProps = new Properties();
        kafkaProps.put("bootstrap.servers", "YOURAWSIP:9092");
        kafkaProps.put("key.serializer", "org.apache.kafka.common.serialization.StringSerializer");
        kafkaProps.put("value.serializer", "org.apache.kafka.common.serialization.StringSerializer");
        KafkaProducer<String, String> producer = new KafkaProducer<String, String>(kafkaProps);

        //Fire-and-forget
        System.out.println("Fire-and-forget starting...");
        ProducerRecord<String, String> record = new ProducerRecord<>("EXERCISE3", "Precision
Products", "France");
        try
        { producer.send(record); }
        catch (Exception e) { e.printStackTrace();}
        System.out.println("Fire-and-forget stopped.");

        //Synchronous send
        System.out.println("Synchronous send starting...");
        ProducerRecord<String, String> record2 =
        new ProducerRecord<>("EXERCISE3", "Precision Products", "France2");
        try { System.out.println("Result get =" + producer.send(record2).get()); }
        catch (Exception e) {e.printStackTrace(); }
        System.out.println("Synchronous stopped.");
    }
}

```

7. Compile and try your project.

C. Create a JAVA Kafka Consumer (for windows): example of

The goal of this section is to create a micro service in JAVA to consume messages from a Kafka service that has already been provisioned, a topic has been created, and messages are being sent to the topic.

1. Create a new Maven project (accordingly with section B).
2. In Project > Properties > Java compiler check that compilation is done using java 8
3. Define the following Maven dependency in pom.xml of your project:

```
<dependencies>
<dependency>
  <groupId>org.apache.kafka</groupId>
  <artifactId>kafka-clients</artifactId>
  <version>2.1.0</version>
</dependency>
</dependencies>
```

4. Create a new Java class (accordingly with section B)
5. Study the following source code and adapt it to your new class:

```
import java.util.Collections;
import java.util.Properties;

import org.apache.kafka.clients.consumer.ConsumerRecord;
import org.apache.kafka.clients.consumer.ConsumerRecords;
import org.apache.kafka.clients.consumer.KafkaConsumer;

public class ClientConsumerKafka {

    public static void main(String[] args) {
        // TODO Auto-generated method stub

        Properties props = new Properties();
        props.put("bootstrap.servers", "yourAWSIP:9092");
        props.put("group.id", "CountryCounter");
        props.put("key.deserializer", "org.apache.kafka.common.serialization.StringDeserializer");
        props.put("value.deserializer", "org.apache.kafka.common.serialization.StringDeserializer");
        KafkaConsumer<String, String> consumer = new KafkaConsumer<String, String>(props);
        consumer.subscribe(Collections.singletonList("EXERCISE3"));
        try {
            while (true)
            {
                /*
                consumers must keep polling Kafka or they will be considered
                dead and the partitions they are consuming will be handed to another
                consumer in the group to continue consuming. The parameter we pass, poll(),
                is a timeout interval and controls how long poll() will block if data is not available
                in the consumer buffer. If this is set to 0, poll() will return immediately;
                otherwise, it will wait for the specified number of milliseconds for data to arrive
                from the broker.

                */
                ConsumerRecords<String, String> records = consumer.poll(100);
                for (ConsumerRecord<String, String> record : records)
                {
                    System.out.printf("topic = %s, partition = %s, offset = %d, customer = %s, country = %s\n",
                        record.topic(), record.partition(), record.offset(),
                        record.key(), record.value());
                }
            }
            finally {consumer.close();}
        }
    }
}
```

6. Compile and try your project.

D. Producing and Consuming an XML message (for windows): example of

The goal of this section is to present how to send an XML message to your consumer.

1. Start the Eclipse Consumer as explained before in this tutorial.
2. Produce the following XML message in docker client

```
<?xml version="1.0" encoding="UTF-8" ?><!-- This is an XML document example --><invoice xmlns="http://example.org" nr="220706" date="2015-07-22T11:05"><seller tax-id="500600"><name>Restaurant</name></seller><buyer tax-id="555111"><name /></buyer><totals currency="EUR"><value>100</value><tax>30</tax></totals><items><item><value>8</value><desc>Meal</desc></item><item><value>2</value><desc>Water</desc></item></items></invoice>
```

3. Verify the message received in the consumer.

E. Create an AWS RDS Database, accessing it, and inserting data: example of

The goal of this section is to exemplify how to create your own cloud database in AWS, to access its contents through a database workbench, and then creating a JAVA microservice to consume Kafka messages that are then inserted in the cloud database.

1. Go to your AWS account console and click on Database RDS -> My SQL and select the following configurations in step 2:

The screenshot shows the 'Specify DB details' page in the AWS Management Console. The left sidebar contains the 'Amazon RDS' menu with options like Dashboard, Databases, Query Editor, Performance Insights, Snapshots, Automated backups, Reserved instances, Subnet groups, Parameter groups, Option groups, Events, Event subscriptions, and Recommendations. The main content area is titled 'Specify DB details' and shows the configuration steps. Step 2 is 'Specify DB details', and Step 3 is 'Configure advanced settings'. The 'Instance specifications' section includes: DB engine (MySQL Community Edition), License model (general-public-license), DB engine version (MySQL 5.6.40), DB instance class (db.t2.micro — 1 vCPU, 1 GiB RAM), Multi-AZ deployment (No), Storage type (General Purpose (SSD)), and Allocated storage (20 GiB). The 'Settings' section includes: DB instance identifier (mytestdb), Master username (storemessages), Master password (masked with dots), and Confirm password (masked with dots).

Specify DB details

Step 2: Specify DB details

Step 3: Configure advanced settings

Instance specifications

Estimate your monthly costs for the DB instance using the [AWS Simple Monthly Calculator](#)

DB engine: MySQL Community Edition

License model: [Info](#)
general-public-license

DB engine version: [Info](#)
MySQL 5.6.40

Known Issues/Limitations
Review the [Known Issues/Limitations](#) to learn about potential compatibility issues with specific database versions.

Free tier
The Amazon RDS Free Tier provides a single db.t2.micro Instance as well as up to 20 GiB of storage, allowing new AWS customers to gain hands-on experience with Amazon RDS. Learn more about the RDS Free Tier and the Instance restrictions [here](#).
☒ Only enable options eligible for RDS Free Usage Tier [Info](#)

DB instance class: [Info](#)
db.t2.micro — 1 vCPU, 1 GiB RAM

Multi-AZ deployment: [Info](#)
☒ Create replica in different zone
Creates a replica in a different Availability Zone (AZ) to provide data redundancy, eliminate I/O freezes, and minimize latency spikes during system backups.
☐ No

Storage type: [Info](#)
General Purpose (SSD)

Allocated storage: 20 GiB
(Minimum: 20 GiB, Maximum: 20 GiB) Higher allocated storage may improve IOPS performance.

Settings

DB instance identifier: [Info](#)
Specify a name that is unique for all DB instances owned by your AWS account in the current region.
mytestdb

DB instance identifier is case insensitive, but stored as all lower-case, as in "mydbinstance". Must contain from 1 to 63 alphanumeric characters or hyphens (1 to 15 for SQL Server). First character must be a letter. Cannot end with a hyphen or contain two consecutive hyphens.

Master username: [Info](#)
Specify an alphanumeric string that defines the login ID for the master user.
storemessages

Master Username must start with a letter. Must contain 1 to 16 alphanumeric characters.

Master password: [Info](#)
Confirm password: [Info](#)

2. Check if the database is started.
3. Use your MySQL database client (for example MySQL workbench), connect to your RDS AWS database and execute the following scripts:

```
DROP DATABASE IF EXISTS collectedMSGs;
CREATE DATABASE IF NOT EXISTS collectedMSGs;
```

```
USE collectedMSGs;
DROP TABLE IF EXISTS Message;

CREATE TABLE Message( offset INTEGER PRIMARY KEY,
                        groupId VARCHAR(100) NOT NULL,
                        contentKey VARCHAR(100),
                        contentValue VARCHAR(1000));
```

4. Create a new Maven Project in Eclipse.
5. Download the mysql/J connector from <https://dev.mysql.com/downloads/connector/j/> and add it to the external library of your project.
6. Use the following dependency in pom.xml

```
<dependency>
  <groupId>mysql</groupId>
  <artifactId>mysql-connector-java</artifactId>
  <version>8.0.13</version>
</dependency>
```

7. Study the following source code and adapt it to your new class:

```

import java.sql.Connection;
import java.sql.DriverManager;
import java.sql.PreparedStatement;
import java.sql.SQLException;
import java.sql.Statement;
import java.util.ArrayList;
import java.util.Collections;
import java.util.List;
import java.util.Properties;

import org.apache.kafka.clients.consumer.CommitFailedException;
import org.apache.kafka.clients.consumer.ConsumerRecord;
import org.apache.kafka.clients.consumer.ConsumerRecords;
import org.apache.kafka.clients.consumer.KafkaConsumer;
import org.apache.kafka.common.TopicPartition;

public class ClientConsumerBDCloud {

    public static void main(String[] args) {
        // TODO Auto-generated method stub

        Properties props = new Properties();
        props.put("bootstrap.servers", "yourAWSIP:9092");
        props.put("group.id", "lCountryCounter");
        props.put("key.deserializer", "org.apache.kafka.common.serialization.StringDeserializer");
        props.put("value.deserializer", "org.apache.kafka.common.serialization.StringDeserializer");
        props.put("auto.commit.offset", "false"); //to commit manually
        KafkaConsumer<String, String> consumer = new KafkaConsumer<String, String>(props);
        consumer.subscribe(Collections.singletonList("EXERCISE3"));
        Connection conn = null;
        boolean bd_ok = false;
        try {
            Class.forName("com.mysql.cj.jdbc.Driver");
            conn = DriverManager.getConnection("jdbc:mysql://yourAWSDBIP:3306/YOURDATABASENAME",
                                                "YOURMasterUSERNAME", "YOURPASSWORD");

            bd_ok = true;
        }
        catch (SQLException sqle) { System.out.println("SQLException: " + sqle); }
        catch (ClassNotFoundException e) { e.printStackTrace(); }

        try {
            while (true)
            {
                ConsumerRecords<String, String> records = consumer.poll(100);
                for (ConsumerRecord<String, String> record : records)
                {
                    System.out.printf("topic = %s, partition = %s, offset = %d,customer = %s, country = %s\n",record.topic(), record.partition(), record.offset(), record.key(), record.value());

                    if (bd_ok)
                    {
                        PreparedStatement s;
                        s = conn.prepareStatement ("insert into Message values(?,?,?,?)");
                        s.setInt(1, new Long(record.offset()).intValue());
                        s.setString(2, props.getProperty("group.id"));
                        s.setString(3, record.key());
                        s.setString(4, record.value());
                        s.executeUpdate();
                        s.close();
                    }

                }
                //Commit Current Offset
                try { consumer.commitSync(); }
                catch (CommitFailedException e) { System.out.printf("commit failed", e);}
            }
        } catch (SQLException e) { e.printStackTrace(); }
        finally {
            consumer.close();
            try {
                conn.close();
            } catch (SQLException e) { e.printStackTrace(); }
        }
    }
}

```

8. Compile and try your project.

F. Deploying JAVA services in EC2: example of

1. Create an EC2 Amazon linux image instance
2. Check that JDK 8 is available in the new created instance with the following commands:

```
[ec2-user@ip- - - ~]$ java -version
openjdk version "1.8.0_191"
OpenJDK Runtime Environment (build 1.8.0_191-b12)
OpenJDK 64-Bit Server VM (build 25.191-b12, mixed mode)
[ec2-user@ip- - - ~]$ javac -version
javac 1.8.0_191
```

Hint 1: if not available, use the following commands:

```
[ec2-user@ip- - - ~]$ yum search jdk
Loaded plugins: priorities, update-motd, upgrade-helper
===== N/S matched: jdk =====
copy-jdk-configs.noarch : JDKs configuration files copier
java-1.6.0-openjdk.x86_64 : OpenJDK Runtime Environment
java-1.6.0-openjdk-demo.x86_64 : OpenJDK Demos
java-1.6.0-openjdk-devel.x86_64 : OpenJDK Development Environment
java-1.6.0-openjdk-javadoc.x86_64 : OpenJDK API Documentation
java-1.6.0-openjdk-src.x86_64 : OpenJDK Source Bundle
java-1.7.0-openjdk.x86_64 : OpenJDK Runtime Environment
java-1.7.0-openjdk-demo.x86_64 : OpenJDK Demos
java-1.7.0-openjdk-devel.x86_64 : OpenJDK Development Environment
java-1.7.0-openjdk-javadoc.noarch : OpenJDK API Documentation
java-1.7.0-openjdk-src.x86_64 : OpenJDK Source Bundle
java-1.8.0-openjdk.x86_64 : OpenJDK Runtime Environment
java-1.8.0-openjdk-demo.x86_64 : OpenJDK Demos
java-1.8.0-openjdk-devel.x86_64 : OpenJDK Development Environment
java-1.8.0-openjdk-headless.x86_64 : OpenJDK Runtime Environment
java-1.8.0-openjdk-javadoc.noarch : OpenJDK API Documentation
java-1.8.0-openjdk-javadoc-zip.noarch : OpenJDK API Documentation compressed in
                                     : single archive
java-1.8.0-openjdk-src.x86_64 : OpenJDK Source Bundle
ldapjdk-javadoc.noarch : Javadoc for ldapjdk
ldapjdk.noarch : The Mozilla LDAP Java SDK

[ec2-user@ip- - - ~]$ sudo yum install java-1.8.0-openjdk-devel.x86_64
Loaded plugins: priorities, update-motd, upgrade-helper
amzn-main | 2.1 kB 00:00
amzn-updates | 2.5 kB 00:00
Package 1:java-1.8.0-openjdk-devel-1.8.0.191.b12-0.42.amzn1.x86_64 already installed and
latest version
Nothing to do
```

Upload apache-maven binary file that is downloaded from URL referred in the end of this document (section Appendix: URLs with other resources).

3. Execute the following commands in the EC2 instance:

```
[ec2-user@ip- - - ~]$ tar xzvf apache-maven-3.6.0-bin.tar.gz

[ec2-user@ip- - - ~]$ cd apache-maven-3.6.0/bin

[ec2-user@ip- - - bin]$ pwd
/home/ec2-user/apache-maven-3.6.0/bin

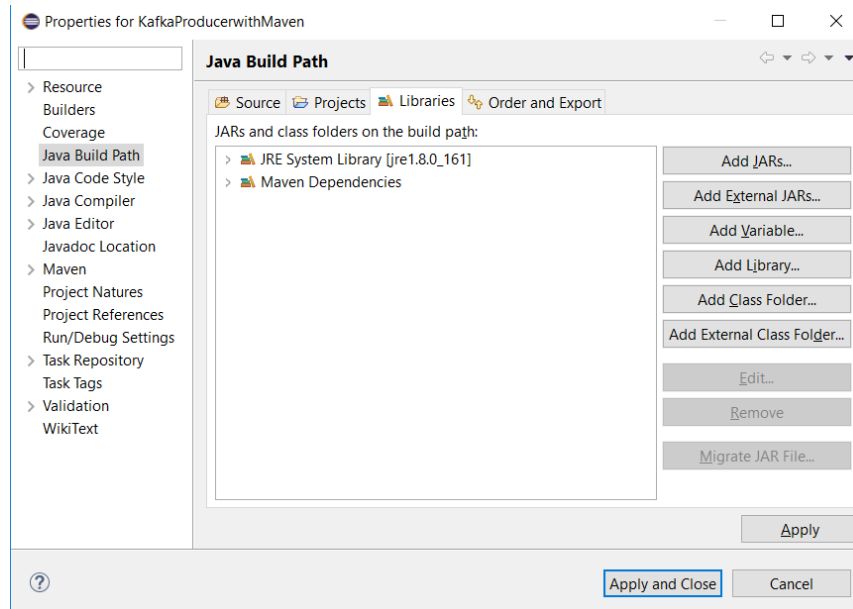
[ec2-user@ip- - - bin]$ export PATH=/home/ec2-user/apache-maven-3.6.0/bin:$PATH

[ec2-user@ip- - - ~]$ mvn -version
Apache Maven 3.6.0 (97c98ec64a1fdfee7767ce5fffb20918da4f719f3; 2018-10-24T18:41:47Z)
Maven home: /home/ec2-user/apache-maven-3.6.0
Java version: 1.8.0_191, vendor: Oracle Corporation, runtime: /usr/lib/jvm/java-1.8.0-openjdk-1.8.0.191.b12-0.42.amzn1.x86_64/jre
Default locale: en_US, platform encoding: UTF-8
OS name: "linux", version: "4.14.77-70.59.amzn1.x86_64", arch: "amd64", family: "unix"
```

Hint 2: To automate the maven path with instance login update the export command in the file:

```
.bash_profile
```

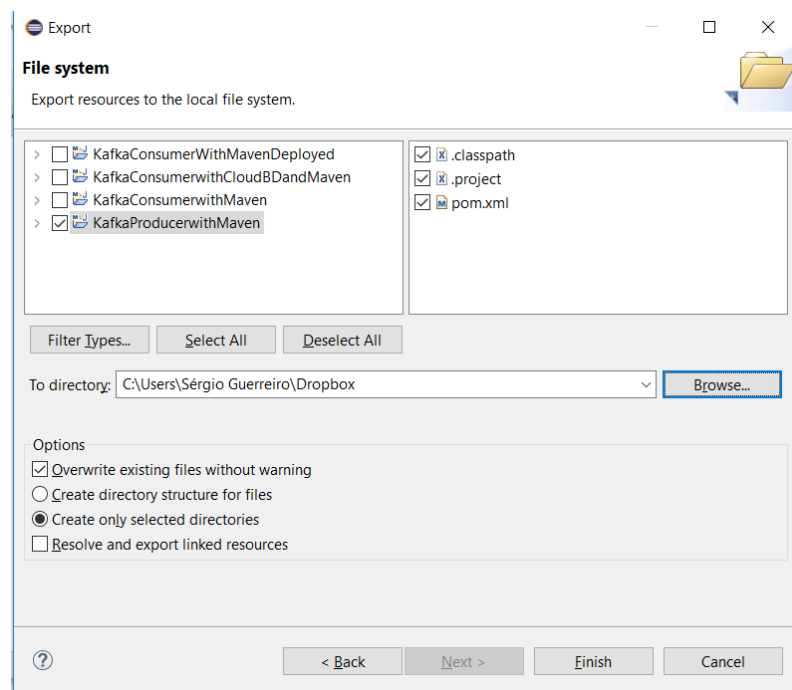
4. Check in Project -> Properties -> Add Library that JAVA8 is added to the environment



5. Also, in pom.xml file add the following:

```
<properties>
  <maven.compiler.source>1.8</maven.compiler.source>
  <maven.compiler.target>1.8</maven.compiler.target>
</properties>
```

6. Right click in the Eclipse project explorer, then Export project as a file system from Eclipse to your computer, as follows:



7. Upload the exported folder project to EC2 AMI using FileZilla.
8. Execute the following commands inside the exported folder project:

```
[ec2-user@ip- - - KafkaProducerwithMaven]$ mvn exec:java -Dexec.mainClass="ClientProducerKafka"
[INFO] Scanning for projects...
[INFO] -----< com.ie.microservices:KafkaProducerwithMaven >-----
[INFO] Building My project 0.0.1-SNAPSHOT
[INFO] -----[ jar ]-----
[INFO] --- exec-maven-plugin:1.6.0:java (default-cli) @ KafkaProducerwithMaven ---
SLF4J: Failed to load class "org.slf4j.impl.StaticLoggerBinder".
SLF4J: Defaulting to no-operation (NOP) logger implementation
```

```
SLF4J: See http://www.slf4j.org/codes.html#StaticLoggerBinder for further details.  
Fire-and-forget starting...  
Fire-and-forget stopped.  
Synchronous send starting...  
Result get =EXERCISE3-0@13  
Synchronous stopped.  
Asynchronous send starting...  
Asynchronous send stopped.
```

G. Dreamfactory installation and operation in AWS EC2

1. Download the version of linux operating system from <https://www.dreamfactory.com/downloads>
2. Create AWS EC2 instance – AMI and start instance.
3. SFTP the downloaded file in step 1. to your EC2 new instance.
4. If your instance has less than 1Gb of RAM, execute the following sequence of commands. Otherwise, jump directly to the command “sudo ./bitnami-dreamfactory-2.14.1-4-linux-x64-installer.run”.

```
Authenticating with public key "imported-openssh-key"
Last login: Tue Jan 15 12:03:41 2019 from kdbio47.inesc-id.pt

  _ | _ | _ | _ |
  | | ( | _ /
  _ | \ _ | _ |
Amazon Linux AMI

https://aws.amazon.com/amazon-linux-ami/2018.03-release-notes/

[ec2-user@ip-172-31-85-254 ~]$ sudo yum update

[ec2-user@ip-172-31-85-254 ~]$ free -m
              total        used         free      shared    buffers     cached
Mem:           985         715          270           0         16         615
-/+ buffers/cache:          83          902
Swap:           0           0           0

[ec2-user@ip-172-31-85-254 /]$ sudo dd if=/dev/zero of=/mnt/swap.0 bs=1024 count=1048576
1048576+0 records in
1048576+0 records out
1073741824 bytes (1.1 GB) copied, 14.0105 s, 76.6 MB/s

[ec2-user@ip-172-31-85-254 /]$ sudo chmod 600 /mnt/swap.0

[ec2-user@ip-172-31-85-254 /]$ sudo mkswap /mnt/swap.0
Setting up swapspace version 1, size = 1048572 KiB
no label, UUID=ba957b3b-8175-41b2-8441-d362e082b9a7

[ec2-user@ip-172-31-85-254 /]$ sudo swapon /mnt/swap.0

[ec2-user@ip-172-31-85-254 /]$ sudo swapon -s
Filename                                Type           Size      Used     Priority
/mnt/swap.0                            file          1048572  0        -2

[ec2-user@ip-172-31-85-254 /]$ cd

[ec2-user@ip-172-31-85-254 ~]$ sudo ./bitnami-dreamfactory-2.14.1-4-linux-x64-installer.run
-----
Welcome to the Bitnami DreamFactory Stack Setup Wizard.

-----
Select the components you want to install; clear the components you do not want
to install. Click Next when you are ready to continue.

NGINX Open Source [Y/n] :Y

DreamFactory : Y (Cannot be edited)

PhpMyAdmin [Y/n] :Y

Is the selection above correct? [Y/n]: Y

-----
Installation folder

Please, choose a folder to install Bitnami DreamFactory Stack

Select a folder [/opt/dreamfactory-2.14.1-4]:

-----
Please enter your database root user password

MariaDB and MongoDB root password :
Re-enter password :
Do you want to configure mail support? [y/N]: N
```

```

-----
Setup is now ready to begin installing Bitnami DreamFactory Stack on your
computer.

Do you want to continue? [Y/n]: Y

-----
Please wait while Setup installs Bitnami DreamFactory Stack on your computer.

Installing
0% _____ 50% _____ 100%
#####

-----
Setup has finished installing Bitnami DreamFactory Stack on your computer.

Launch Bitnami DreamFactory Stack [Y/n]: Y

[ec2-user@ip-172-31-85-254 ~]$ free -m
              total        used        free      shared    buffers     cached
Mem:           985         620         365          18          53          321
-/+ buffers/cache:         245         740
Swap:          1023           0         1023

[ec2-user@ip-172-31-85-254 ~]$ cd /opt/dreamfactory-2.14.1-4/

[ec2-user@ip-172-31-85-254 dreamfactory-2.14.1-4]$ sudo ./ctlscript.sh status
apache already running
mysql already running
mongodb already running

```

Hint1: after an instance reboot, the dreamfactory is restarted with:

```

Using username "ec2-user".
Authenticating with public key "imported-openssh-key"
Last login: Tue Jan 15 12:07:21 2019 from kdbio47.inesc-id.pt

  _ | _ | _ )
  _ | ( _ /   Amazon Linux AMI
  _ | \ _ | _ |

https://aws.amazon.com/amazon-linux-ami/2018.03-release-notes/
[ec2-user@ip-172-31-85-254 ~]$ cd /opt/dreamfactory-2.14.1-4/

[ec2-user@ip-172-31-85-254 dreamfactory-2.14.1-4]$ sudo ./ctlscript.sh status
apache not running
mysql not running
mongodb not running

[ec2-user@ip-172-31-85-254 dreamfactory-2.14.1-4]$ sudo ./ctlscript.sh start
/opt/dreamfactory-2.14.1-4/mysql/scripts/ctl.sh : mysql  started at port 3306
/opt/dreamfactory-2.14.1-4/mongodb/scripts/ctl.sh : mongodb started at port 27017
Syntax OK
/opt/dreamfactory-2.14.1-4/apache2/scripts/ctl.sh : httpd started at port 80

[ec2-user@ip-172-31-85-254 dreamfactory-2.14.1-4]$ sudo ./ctlscript.sh status
apache already running
mysql already running
mongodb already running

```

Your dreamfactory is now installed and running. To test the invocation of the REST interface, execute the following steps:

5. Access your dreamfactory server at <http://YourAWSInstanceName:8080>
6. Create your account
7. Choose Services -> Create button -> Select Service Type -> Database -> MySQL
8. Use the following configuration in Info tab:
 - Name:* TestMySQLAPI
 - Label:* for testing of TestMySQLAPI
 - Description:* this setup is for testing of TestMySQLAPI
9. Use the following configuration in Config tab:

Host: your Amazon RDS endpoint. It can be consulted at AWS Console -> RDS -> Databases

Port Number: 3306 (default)

Database: the name of your database previously created

Username: your username

Password: your password

10. Press Save button

11. *Optional:* You can check a set of interactive Swagger documentation, in API Docs

12. Secure your REST API, in Roles-> Create Button

13. Use the following configuration in Basic tab:

Name: aws-mysql

Description: this setup is for testing of TestMySQLAPI

Check Active

14. Use the following configuration in Access tab:


Press "Add new service access" in the "+" button, and choose:

Service: your service defined in step 8.

Component: your table name, e.g., "_table/Message/*"

Access: choose the access that you would like to grant from - GET, POST, PUT, PATCH, DELETE

15. Press Save button

16. Create a new API key and associate the key with this role, in  -> Create Button, and use the following configuration.

Application Name: MySQLAWS


Description: The name of your description

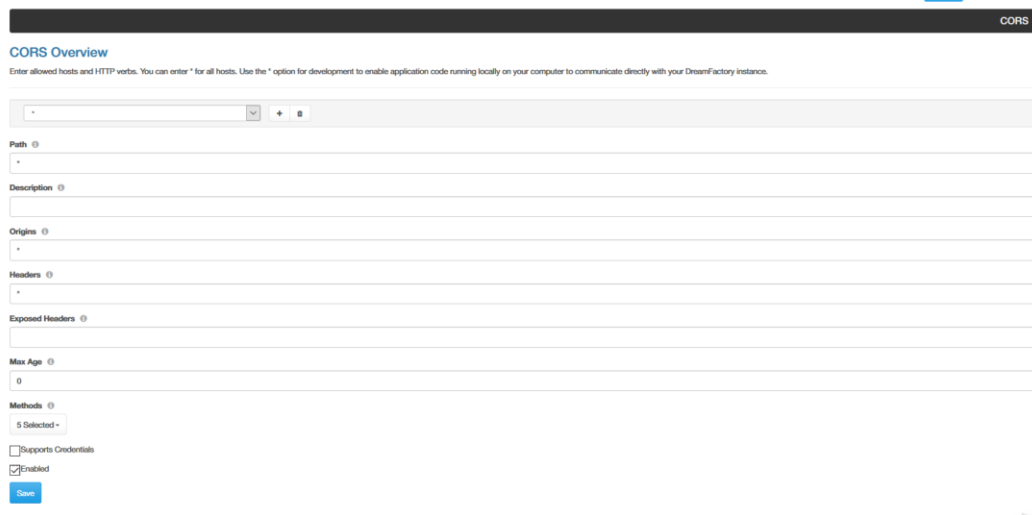
Assign a Default Role: your role name defined in step 13

App Location: No Storage Required - remote device, client, or desktop.

Check Active

17. Press the Save button and you'll be returned to the Apps index screen where the new API key can be copied!
Copy the key into a text file for later reference.

18. Enable CORS (Cross-Origin Resource Sharing) for the new API, in  -> CORS. It will allow API-restricted traffic from all network addresses. Use the following configuration.



19. Test your REST API using the URLs:

http://YourAWSIP:8080/api/v2/YourService/_table/YourTable?api_key=yourAPIKey

Hint 2: If you need to filter results then use: https://{url}/api/v2/{api_name}/_table/{table_name}?filter={filter_string}, e.g.: https://{url}/api/v2/{api_name}/_table/{table_name}?api_key=yourAPIKEY&filter=groupId%20like%201CountryCounter3

```
Pay Attention for "\" ->
curl -v
URL:8080/api/v2/TestMySQLAPIable/_table/Message?api_key=apikey\&filter=groupId%20like%201CountryCounter
```

or using a CURL command:

```
PS C:\Users\Sérgio Guerreiro> curl -v
35.180.252.107:8080/api/v2/TestMySQLAPI/_table/Message?api_key=e114bbf45781ad62a70
dee9a16db26a86a852c2f8c44ec858699d20f46c1f7d9
VERBOSE: GET
http://35.180.252.107:8080/api/v2/TestMySQLAPI/_table/Message?api_key=e114bbf45781ad62a70dee9a16db26a86a852c2
f8c44ec858
699d20f46c1f7d9 with 0-byte payload
VERBOSE: received 403-byte response of content type application/json

StatusCode      : 200
StatusDescription : OK
Content         :
{"resource":[{"offset":182,"groupId":"1CountryCounter","contentKey":null,"contentValue":"sdssffd"},
               {"offset":183,"groupId":"1CountryCounter","contentKey":"Precision
               Products","contentValue":"France"},...
RawContent      : HTTP/1.1 200 OK
                  X-Frame-Options: SAMEORIGIN
                  Vary: Cookie
                  Keep-Alive: timeout=5, max=100
                  Connection: Keep-Alive
                  Content-Length: 403
                  Cache-Control: no-cache, private
                  Content-Type: application/jso...
Forms           : {}
Headers         : {[X-Frame-Options, SAMEORIGIN], [Vary, Cookie], [Keep-Alive, timeout=5, max=100],
[Connection,
                  Keep-Alive]...}
Images          : {}
InputFields     : {}
Links           : {}
ParsedHtml      : mshtml.HTMLDocumentClass
RawContentLength : 403
```

Hint 3: automate your REST API tests using Postman tool.

H. Creating a topic in Kafka from JAVA

Follow the similar steps presented in section "B. Create a JAVA Kafka Producer (for windows): example of" of this document with the following solution:

Maven dependencies:

```
<dependency>
  <groupId>org.apache.kafka</groupId>
  <artifactId>kafka 2.11</artifactId>
  <version>0.10.2.0</version>
</dependency>
<dependency>
  <groupId>com.ioItec</groupId>
  <artifactId>zkclient</artifactId>
  <version>0.9</version>
</dependency>
```

Java code:

```
import org.ioItec.zkclient.ZkClient;
import org.ioItec.zkclient.ZkConnection;

import java.util.Properties;

import kafka.admin.AdminUtils;
import kafka.admin.RackAwareMode;
import kafka.utils.ZKStringSerializer$;
import kafka.utils.ZkUtils;

public class KafkaJavaExample {

    public static void main(String[] args) {
        String zookeeperConnect = "zkserver1:2181";
        int sessionTimeoutMs = 10 * 1000;
        int connectionTimeoutMs = 8 * 1000;

        String topic = "my-NEW-topic";
        int partitions = 1;
        int replication = 1;
        Properties topicConfig = new Properties(); // add per-topic configurations settings here

        // Note: You must initialize the ZkClient with ZKStringSerializer. If you don't, then
        // createTopic() will only seem to work (it will return without error). The topic will exist in
        // only ZooKeeper and will be returned when listing topics, but Kafka itself does not create the
        // topic.
        ZkClient zkClient = new ZkClient(
            zookeeperConnect,
            sessionTimeoutMs,
            connectionTimeoutMs,
            ZKStringSerializer$.MODULE$);

        // Security for Kafka was added in Kafka 0.9.0.0
        boolean isSecureKafkaCluster = false;

        ZkUtils zkUtils = new ZkUtils(zkClient, new ZkConnection(zookeeperConnect), isSecureKafkaCluster);
        AdminUtils.createTopic(zkUtils, topic, partitions, replication, topicConfig,
            RackAwareMode.Enforced$.MODULE$);
        zkClient.close();
    }
}
```

I. Appendix: URLs with other resources

- To create a Maven project in eclipse: <https://www.tech-recipes.com/rx/39279/create-a-new-maven-project-in-eclipse/>
- About Kafka group-ids: <https://stackoverflow.com/questions/35561110/can-multiple-kafka-consumers-read-same-message-from-the-partition>
- More details about DreamFactory usage: <https://blog.dreamfactory.com/create-a-mysql-rest-api-in-minutes-using-dreamfactory/>
- Where to download apache maven <https://maven.apache.org/download.cgi>
- How to install apache maven <https://maven.apache.org/install.html>
- Postman tool: <https://www.getpostman.com/postman>
- http://wiki.dreamfactory.com/DreamFactory/Tutorials/Querying_records_with_logical_filters